

Polar

October 03, 2023

1 Imports

```
[1]: import json
      from functools import reduce
      from os import listdir
      from os.path import isfile, join
      from pathlib import Path

      import matplotlib.pyplot as plt
      import numpy as np
      import pandas as pd
      import statsmodels.api as sm
      import statsmodels.formula.api as smf
      import tabulate
      from pylab import rcParams
      from scipy.stats import shapiro
      from statsmodels.graphics.gofplots import qqplot
      from statsmodels.stats.diagnostic import het_goldfeldquandt
      from statsmodels.stats.outliers_influence import variance_inflation_factor

[2]: from IPython.display import display, HTML
```

2 Setup

```
[3]: # Set figure size
      rcParams["figure.figsize"] = (4, 4)

      # Folder for images
      Path("img").mkdir(parents=True, exist_ok=True)

      # Nice float format
      pd.options.display.float_format = "{:,.2f}".format
```

3 Data description

Last year I purchased a Polar watch that tracks my vitals during workouts. I used the [Polar Flow](#) website to obtain a copy of my data. For privacy reasons I shall not be sharing the dataset.

```
[4]: path = './data/'
```

First, we create a list of files in the download.

```
[5]: files = [f for f in listdir(path) if isfile(join(path, f))]
```

We shall only consider files containing the string 'training-session'.

```
[6]: files = [f for f in files if 'training-session' in f]
```

The number of files under consideration is:

```
[7]: len(files)
```

```
[7]: 284
```

We loop over each of the files and them to a list.

```
[8]: data = []

for f in files:
    with open(join(path, f)) as f:
        d = json.load(f)
        data.append(d)
```

We define a function to extract statistics about heart rate measured during the workouts.

```
[9]: quantiles = [0.01, 0.25, 0.5, 0.75, 0.99]
```

```
[10]: def extract_hr_info(workout, quantiles):

    stats = {'heartRateAvg2': np.nan,
             'heartRateStd': np.nan}

    for q in quantiles:
        stats[f'heartRateQ' + str(int(q * 100))] = np.nan

    # Check if data exists
    try:
        heart_rates = workout['exercises'][0]['samples']['heartRate']
    except KeyError:
        return stats

    # Loop over measurements
```

```

hr_data = []
for hr in heart_rates:

    # Check if actually measured hr
    if 'value' in hr:
        hr_data.append(hr['value'])

stats['heartRateAvg2'] = np.mean(hr_data)
stats['heartRateStd'] = np.std(hr_data)

for q in quantiles:
    stats[f'heartRateQ' + str(int(q * 100))] = np.quantile(hr_data, q)

return stats

```

We extract the relevant information from the items in the list.

```

[11]: workouts = []

for d in data:
    basic = d['exercises'][0]
    hr = extract_hr_info(workout=d,
                        quantiles=quantiles)

    workouts.append(**basic, **hr)

```

Finally we create a dataframe containing the workout information.

```

[12]: df = pd.DataFrame(workouts)

```

4 Data structure

We find the following columns in the dataframe.

```

[13]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284 entries, 0 to 283
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   startTime              284 non-null    object
1   stopTime               284 non-null    object
2   timezoneOffset         284 non-null    int64
3   duration               284 non-null    object
4   sport                  284 non-null    object
5   kiloCalories           283 non-null    float64
6   heartRate              283 non-null    object

```

```

7  zones                284 non-null    object
8  samples              284 non-null    object
9  heartRateAvg2        283 non-null    float64
10 heartRateStd         283 non-null    float64
11 heartRateQ1          283 non-null    float64
12 heartRateQ25         283 non-null    float64
13 heartRateQ50         283 non-null    float64
14 heartRateQ75         283 non-null    float64
15 heartRateQ99         283 non-null    float64
16 distance             130 non-null    float64
17 latitude             130 non-null    float64
18 longitude            130 non-null    float64
19 ascent               120 non-null    float64
20 descent              121 non-null    float64
21 speed                130 non-null    object
22 autoLaps             102 non-null    object
23 laps                 2 non-null     object

```

dtypes: float64(13), int64(1), object(10)

memory usage: 53.4+ KB

We remove columns that contain data from features I do not use in my training.

Due to privacy concerns I shan't be extracting longitudinal and latitudinal data.

```
[14]: df = df.drop(['zones', 'samples', 'autoLaps',
                  'laps', 'latitude', 'longitude',
                  'ascent', 'descent'], axis=1)
```

```
[15]: df.head()
```

```
[15]:
```

	startTime	stopTime	timezoneOffset	\
0	2019-05-24T13:18:14.000	2019-05-24T14:58:44.125	120	
1	2019-05-04T12:03:34.000	2019-05-04T13:21:38.500	120	
2	2019-04-12T12:48:57.000	2019-04-12T12:59:10.750	120	
3	2019-06-12T13:13:09.000	2019-06-12T13:23:15.500	120	
4	2019-05-24T14:59:06.000	2019-05-24T15:29:08.750	120	

	duration	sport	kiloCalories	\
0	PT6030.125S	STRENGTH_TRAINING	658.00	
1	PT4684.500S	STRENGTH_TRAINING	373.00	
2	PT613.750S	TREADMILL_RUNNING	62.00	
3	PT606.500S	TREADMILL_RUNNING	71.00	
4	PT1802.750S	TREADMILL_RUNNING	416.00	

	heartRate	heartRateAvg2	heartRateStd	\
0	{'min': 72, 'avg': 105, 'max': 136}	104.77	11.28	
1	{'min': 71, 'avg': 99, 'max': 138}	98.65	12.51	
2	{'min': 71, 'avg': 97, 'max': 107}	97.07	8.00	

3	{'min': 67, 'avg': 105, 'max': 121}	105.24	11.25
4	{'min': 84, 'avg': 144, 'max': 170}	143.85	18.47

	heartRateQ1	heartRateQ25	heartRateQ50	heartRateQ75	heartRateQ99 \
0	77.00	99.00	105.00	111.00	132.00
1	74.00	91.00	97.00	106.00	126.00
2	72.00	94.00	97.00	104.00	107.00
3	67.96	98.00	104.00	118.00	121.00
4	87.00	133.00	146.00	158.00	169.00

	distance	speed
0	nan	NaN
1	nan	NaN
2	nan	NaN
3	nan	NaN
4	nan	NaN

5 Missing Values

The watch tracks different information for different workouts. For example when walking it tracks location but when walking on a treadmill it doesn't, hence there is quite a lot of missing data.

```
[16]: missing = (df.isna().sum() / df.shape[0] * 100)
missing.name = 'Percent missing'
missing = missing.to_frame()
missing = missing.sort_values('Percent missing', ascending=False)
missing = missing[missing['Percent missing'] > 0]
missing = missing.reset_index()
missing = missing.rename(columns={'index': 'Feature'})
np.round(missing, 2)
```

```
[16]:
```

	Feature	Percent missing
0	distance	54.23
1	speed	54.23
2	kiloCalories	0.35
3	heartRate	0.35
4	heartRateAvg2	0.35
5	heartRateStd	0.35
6	heartRateQ1	0.35
7	heartRateQ25	0.35
8	heartRateQ50	0.35
9	heartRateQ75	0.35
10	heartRateQ99	0.35

6 Transforms

We apply certain transforms to make the data easier to work with. First we convert strings to datetimes.

```
[17]: df["startTime"] = pd.to_datetime(df["startTime"])
      df["stopTime"] = pd.to_datetime(df["stopTime"])
```

We calculate the total duration of each individual workout in minutes.

```
[18]: df["totalTime"] = (df["stopTime"] - df["startTime"])
      df["totalTime"] = df["totalTime"].apply(lambda x: round(x.seconds / 60, 2))
      df.drop("duration", axis=1, inplace=True)
```

We extract maximum, average and minimum heart rate values from the heartRate column.

```
[19]: df["heartRateMax"] = df["heartRate"].apply(lambda x: x["max"] if isinstance(x, dict) else np.nan)
      df["heartRateAvg"] = df["heartRate"].apply(lambda x: x["avg"] if isinstance(x, dict) else np.nan)
      df["heartRateMin"] = df["heartRate"].apply(lambda x: x["min"] if isinstance(x, dict) else np.nan)
      df.drop("heartRate", axis=1, inplace=True)
```

We assume that if there is no distance then the workout was indoors:

```
[20]: df["isInside"] = df["distance"].apply(lambda x: True if pd.isnull(x) else False)
      df = df.drop(["distance", "speed"], axis=1)
```

We are going to map sports to different activityType's. We will map strength training to 1 and cardiovascular work to 0.

```
[21]: def sport_to_activity_type(x):
      if "strength" in x.lower():
          return True
      else:
          return False
```

```
[22]: df["isStrength"] = df["sport"].apply(sport_to_activity_type)
```

```
[23]: df["sport"] = df["sport"].apply(lambda x: x.lower())
      df["sport"] = pd.Categorical(df["sport"])
```

We extract a list of unique sport values:

```
[24]: sports = sorted(list(df["sport"].unique()))
```

We reorder the alphabetically

```
[25]: order = sorted(df.columns.to_list())
```

```
[26]: df = df[order]
```

We check if there are any more NaN's in the data.

```
[27]: df.isna().sum()
```

```
[27]: heartRateAvg      1
      heartRateAvg2    1
      heartRateMax     1
      heartRateMin     1
      heartRateQ1      1
      heartRateQ25     1
      heartRateQ50     1
      heartRateQ75     1
      heartRateQ99     1
      heartRateStd     1
      isInside         0
      isStrength       0
      kiloCalories     1
      sport            0
      startTime        0
      stopTime         0
      timezoneOffset   0
      totalTime        0
      dtype: int64
```

There is one row with NaN's. This might due to my watch having little battery left to make the measurements.

```
[28]: df = df.dropna()
```

We proceed to sort the data with the latest workouts at the top of the dataframe.

```
[29]: sort_cols = ["startTime", "startTime"]
      df = df.sort_values(sort_cols, ascending=False)
      df = df.reset_index(drop=True)
```

We verify that the datatypes are correct.

```
[30]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 283 entries, 0 to 282
Data columns (total 18 columns):
#   Column          Non-Null Count  Dtype
---  -
0   heartRateAvg     283 non-null   float64
1   heartRateAvg2    283 non-null   float64
2   heartRateMax     283 non-null   float64
```

3	heartRateMin	283	non-null	float64
4	heartRateQ1	283	non-null	float64
5	heartRateQ25	283	non-null	float64
6	heartRateQ50	283	non-null	float64
7	heartRateQ75	283	non-null	float64
8	heartRateQ99	283	non-null	float64
9	heartRateStd	283	non-null	float64
10	isInside	283	non-null	bool
11	isStrength	283	non-null	bool
12	kiloCalories	283	non-null	float64
13	sport	283	non-null	category
14	startTime	283	non-null	datetime64[ns]
15	stopTime	283	non-null	datetime64[ns]
16	timezoneOffset	283	non-null	int64
17	totalTime	283	non-null	float64

dtypes: bool(2), category(1), datetime64[ns](2), float64(12), int64(1)
memory usage: 34.3 KB

[31]: df.head()

```
[31]:
```

	heartRateAvg	heartRateAvg2	heartRateMax	heartRateMin	heartRateQ1	\
0	102.00	102.03	116.00	69.00	73.00	
1	104.00	104.14	132.00	70.00	73.81	
2	103.00	103.10	122.00	87.00	91.00	
3	108.00	107.97	125.00	87.00	91.00	
4	117.00	117.08	141.00	90.00	92.00	

	heartRateQ25	heartRateQ50	heartRateQ75	heartRateQ99	heartRateStd	\
0	96.00	103.00	109.00	115.00	9.02	
1	86.00	110.50	118.00	131.00	16.67	
2	96.00	101.00	110.00	120.00	7.94	
3	103.00	108.00	114.00	124.00	7.46	
4	103.00	120.00	128.00	141.00	13.76	

	isInside	isStrength	kiloCalories	sport	startTime	\
0	False	False	245.00	walking	2020-03-29 21:50:21	
1	False	False	401.00	walking	2020-03-27 20:38:32	
2	False	False	336.00	walking	2020-03-26 21:07:46	
3	False	False	380.00	walking	2020-03-25 19:22:38	
4	False	False	358.00	walking	2020-03-24 13:09:06	

	stopTime	timezoneOffset	totalTime
0	2020-03-29 22:23:41.750	120	33.33
1	2020-03-27 21:25:03.750	60	46.52
2	2020-03-26 21:52:55.625	60	45.15
3	2020-03-25 20:10:17.875	60	47.65
4	2020-03-24 13:48:46.750	60	39.67

7 Data analysis

Given that we have produced a clean dataset we can proceed to analyse a few aspects.

7.1 Time span

The date of the first workout is:

```
[32]: str(df['startTime'].min())
```

```
[32]: '2019-02-20 20:46:35'
```

The date of the last workout is:

```
[33]: str(df['startTime'].max())
```

```
[33]: '2020-03-29 21:50:21'
```

Workouts measured:

```
[34]: len(df)
```

```
[34]: 283
```

7.2 Descriptive statistics

```
[35]: df.drop('timezoneOffset', axis=1).describe()
```

```
[35]:
```

	heartRateAvg	heartRateAvg2	heartRateMax	heartRateMin	heartRateQ1	\
count	283.00	283.00	283.00	283.00	283.00	
mean	105.19	105.24	128.34	76.74	80.61	
std	11.87	11.86	18.25	8.99	8.40	
min	82.00	81.98	93.00	53.00	54.00	
25%	96.00	96.42	115.00	70.00	75.00	
50%	103.00	103.42	125.00	77.00	80.00	
75%	111.00	111.26	138.50	83.00	86.00	
max	148.00	148.35	178.00	99.00	107.00	

	heartRateQ25	heartRateQ50	heartRateQ75	heartRateQ99	heartRateStd	\
count	283.00	283.00	283.00	283.00	283.00	
mean	98.24	105.55	112.64	125.84	10.52	
std	10.69	12.26	14.36	18.06	4.46	
min	77.00	82.00	87.00	92.00	2.96	
25%	91.00	97.00	102.00	113.00	7.58	
50%	97.00	104.00	111.00	123.00	10.00	
75%	103.00	112.00	119.00	135.00	12.25	
max	146.00	151.00	160.00	177.00	27.10	

kiloCalories totalTime

count	283.00	283.00
mean	315.98	42.83
std	218.75	29.65
min	29.00	5.00
25%	121.50	15.92
50%	277.00	36.45
75%	441.50	65.29
max	1,067.00	172.73

7.3 Kilocalories burned in total

First we count the total kiloCalories I burned during the period in question.

```
[36]: total_calories = df["kiloCalories"].sum()
      print(total_calories)
```

89421.0

We convert this number to kilograms of body fat. According to [this article](#) it equates to

```
[37]: def kcal_to_kg(x):
      return round(x / 7700, 2)
```

```
[38]: kcal_to_kg(total_calories)
```

[38]: 11.61

7.4 Kilocalories burned by sport

```
[39]: by_sport = df[["kiloCalories", "sport"]].groupby("sport", as_index=False)
      by_sport = by_sport.sum()
      by_sport["sport"] = by_sport["sport"].apply(lambda x: x.lower())
      by_sport["kiloCalories"] = by_sport["kiloCalories"].astype(int)
      by_sport = by_sport.rename(columns={"kiloCalories": "Total kilocalories",
      ↪ "sport": "Sport"})
      by_sport = by_sport.sort_values("Total kilocalories", ascending=False)
      by_sport["Total kilograms"] = by_sport["Total kilocalories"].apply(kcal_to_kg)

      # by_sport = by_sport.style.background_gradient(cmap='YlGn', subset='Total_
      ↪ kilograms')
      # by_sport = by_sport.set_precision(2)

      by_sport
```

```
[39]:
```

	Sport	Total kilocalories	Total kilograms
4	walking	33080	4.30
2	strength_training	31547	4.10

3	treadmill_running	19825	2.57
0	cycling	4029	0.52
1	running	940	0.12

7.5 Kilocalories burned over time

Next we produce a plot of kiloCalories burned over a two month period in 2019. First we extract the relevant data.

```
[40]: start = pd.to_datetime('2019-04-1')
stop = pd.to_datetime('2019-06-1')

daily = df[['startTime', 'kiloCalories']]
mask = (daily['startTime'] >= start) & (daily['startTime'] < stop)
daily = daily[mask]
daily['startTime'] = daily['startTime'].dt.date
daily = daily.groupby('startTime', as_index=False)
daily = daily.sum()
daily = daily.sort_values('startTime', ascending=False)
daily['startTime'] = pd.to_datetime(daily['startTime'])
daily = daily.reset_index(drop=True)
```

We create a dataframe with all the dates to perform a left join and fill the NaN's with zeroes.

```
[41]: dates = pd.date_range(start, stop)
dates = dates.to_frame()
dates = dates.reset_index(drop=True)
dates.columns = ['startTime']
```

```
[42]: daily = pd.merge(dates, daily, on='startTime', how='left')
daily = daily.fillna(0)
```

Finally we produce the figure:

```
[43]: width = 800
height = 400
dpi = 100

plt.figure(figsize=(width/dpi, height/dpi))
plt.plot(daily['startTime'], daily['kiloCalories'])

plt.fill_between(x=daily['startTime'],
                 y1=0,
                 y2=daily['kiloCalories'],
                 alpha=1/2)


daily_avg = daily['kiloCalories'].mean()
```

```

plt.hlines(xmin=daily["startTime"].min(),
           xmax=daily["startTime"].max(),
           y=daily_avg,
           linestyle='dashed',
           label=f'Daily average = {round(daily_avg)} kcal',
           alpha=1/2)

plt.title('Kilocalories over time')
plt.xticks(rotation=45, horizontalalignment='center')
plt.xlim(daily["startTime"].min(), daily["startTime"].max())
plt.ylim(0, daily["kiloCalories"].max() * 1.05)
plt.ylabel('Kilocalories')
plt.legend(loc='best')
plt.tight_layout()
plt.savefig('./img/kilocalories_ts.png')
plt.show()

```



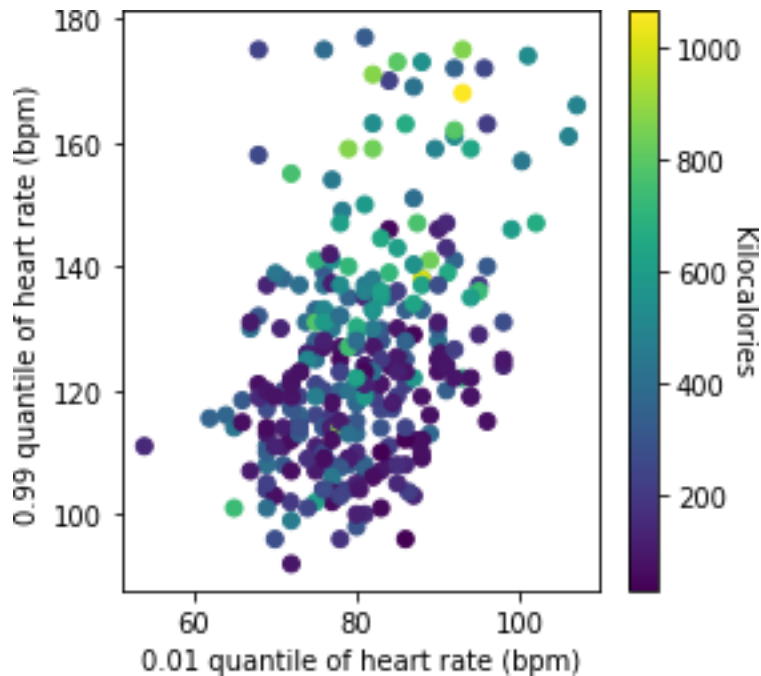
7.6 Kilocalories by intensity

```

[44]: plt.scatter(df["heartRateQ1"], df["heartRateQ99"], c=df["kiloCalories"])
plt.xlabel('0.01 quantile of heart rate (bpm)')
plt.ylabel('0.99 quantile of heart rate (bpm)')

cbar = plt.colorbar()
cbar.set_label('Kilocalories', rotation=270)
plt.savefig('./img/intensity_scatter.png')
plt.show()

```



7.7 Workouts by sport

We check how many workouts I completed.

```
[45]: stats = df[['sport', 'startTime']]
stats = stats.groupby(['sport'], as_index=False)
stats = stats.count()
stats = stats.rename(columns={'sport': 'Sport',
                              'startTime': 'Count'})
stats = stats.sort_values('Count', ascending=False)

# stats = stats.style.background_gradient(cmap='YlGn', subset='Count')
# stats = stats.set_precision(2)

stats
```

```
[45]:
```

	Sport	Count
4	walking	105
3	treadmill_running	90
2	strength_training	62
0	cycling	24
1	running	2

7.8 By hour of day


We count workouts by hour of day.

```
[46]: by_hour = df[["startTime", "sport"]].copy()
      by_hour["startHour"] = by_hour["startTime"].dt.hour
      by_hour = by_hour.drop("startTime", axis=1)
      by_hour = by_hour.groupby("startHour", as_index=False)
      by_hour = by_hour.count()

      all_hours = pd.DataFrame(range(0, 24), columns=["startHour"])

      by_hour = pd.merge(all_hours, by_hour, how="left")
      by_hour = by_hour.fillna(0)
      by_hour = by_hour.sort_values("startHour")
      by_hour = by_hour.rename(columns={"startHour": "Hour of day",
                                       "sport": "Total workouts"})

[47]: plt.bar(by_hour["Hour of day"], by_hour["Total workouts"])
      plt.ylabel("Number of workouts")
      plt.xlabel("Hour of day")
      plt.tight_layout()
      plt.savefig("./img/workouts_by_hour_of_day.png")
      plt.show()
```




7.9 By day of week

We count workouts by day of week.

```
[48]: by_day = df[['startTime', 'sport']].copy()
by_day['Day of week'] = pd.to_datetime(by_day['startTime']).dt.day_name()
by_day['Day number'] = pd.to_datetime(by_day['startTime']).dt.dayofweek
by_day = by_day.groupby(['Day of week', 'Day number'], as_index=False)
by_day = by_day.count()
by_day = by_day.drop('startTime', axis=1)
by_day = by_day.sort_values('Day number')
by_day = by_day.rename(columns={'sport': 'Total Workouts'})
```


```
[49]: plt.bar(by_day['Day of week'], by_day['Total Workouts'])
plt.xticks(rotation=90)
plt.ylabel('Number of workouts')
plt.savefig('./img/workouts_by_day_of_week.png')
plt.show()
```



7.10 Scatter plot of walks data


We plot totalTime versus kiloCalories. As can be seen there seems to exist a linear relationship between the two.

```
[50]: walking = df[df["sport"] == "walking"]
plt.scatter(walking["totalTime"], walking["kiloCalories"], s=2)
plt.xlabel("Duration (minutes)")
plt.ylabel("Kilocalories")
plt.savefig("./img/walks_kilocalories_vs_time.png")
plt.show()
```



We plot heartRateAvg against kiloCalories. Again we see a linear relationship although there are a couple of outliers

```
[51]: walking = df[df["sport"] == "walking"]
plt.scatter(walking["heartRateAvg"], walking["kiloCalories"], s=2)
plt.ylabel("Kilocalories")
plt.xlabel("Average HR (bpm)")
plt.savefig("./img/walks_kilocalories_vs_avg_hr.png")
plt.show()
```

8 Regression

8.1 Data preparation

Now we proceed to build a regression model to predict kiloCalories burned during a workout. First we create a subset of the original data.

```
[52]: reg_df = df[['kiloCalories', 'totalTime',
                  'heartRateQ99', 'isStrength', 'sport']].copy()
```

```
[53]: reg_df.head()
```

```
[53]:   kiloCalories  totalTime  heartRateQ99  isStrength  sport
0         245.00        33.33         115.00         False  walking
1         401.00        46.52         131.00         False  walking
2         336.00        45.15         120.00         False  walking
3         380.00        47.65         124.00         False  walking
4         358.00        39.67         141.00         False  walking
```

We remove the rows where sport is running because there were only two workouts recorded during the period in question.

```
[54]: reg_df = reg_df[reg_df['sport'] != 'running']
```

8.1.1 Outliers

The data is cleansed of outliers using interquartile range.

```
[55]: def is_outlier_iqr(series, k=1.5):  
      """  
      Check if value is an outlier  
      using interquartile range.  
      """  
  
      q1 = series.quantile(0.25)  
      q3 = series.quantile(0.75)  
      iqr = q3 - q1  
      is_outlier = (series <= q1 - k * iqr) | (q3 + k * iqr <= series)  
  
      return is_outlier
```


```
[56]: time_mask = is_outlier_iqr(series=reg_df["totalTime"])  
      kcal_mask = is_outlier_iqr(series=reg_df["kiloCalories"])  
      hr_mask = is_outlier_iqr(series=reg_df["heartRateQ99"])
```

```
[57]: reg_df = reg_df[~(time_mask | kcal_mask | hr_mask)]
```


8.2 Histograms

We proceed to visualize histograms of each of the variables.


```
[58]: plt.hist(reg_df["kiloCalories"], bins=30)  
      plt.xlabel("Kilocalories")  
      plt.ylabel("Frequency")  
      plt.savefig("./img/kilocalories_histogram.png")  
      plt.show()
```



```
[59]: plt.hist(reg_df["totalTime"], bins=30)
plt.xlabel("Duration (minutes)")
plt.ylabel("Frequency")
plt.savefig("./img/duration_histogram.png")
plt.show()
```



```
[60]: plt.hist(reg_df["heartRateQ99"], bins=30)
plt.xlabel("0.99 quantile of heart rate (bpm)")
plt.ylabel("Frequency")
plt.savefig("./img/q99_hr_histogram.png")
plt.show()
```



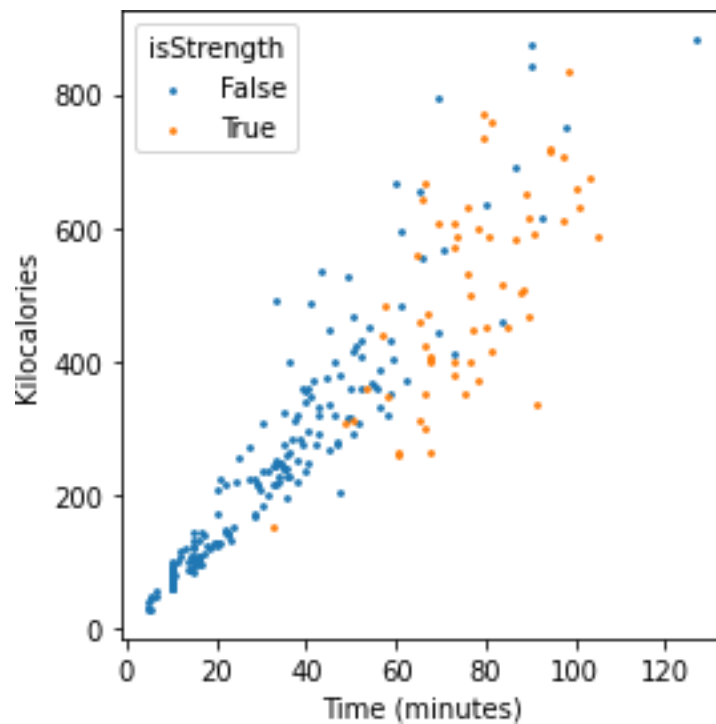
8.3 Scatter plots

The plot below gives reason to suspect a linear relationship between kiloCalories and totalTime.

```
[61]: for val in [False, True]:
    tmp = reg_df[reg_df["isStrength"] == val]
    plt.scatter(tmp["totalTime"],
                tmp["kiloCalories"],
                s=3,
                label=val)


plt.xlabel("Time (minutes)")
plt.ylabel("Kilocalories")
plt.legend(title="isStrength", loc="best")
plt.tight_layout()
plt.savefig("./img/time_vs_kilocalories_scatter_by_strength.png")
```

```
plt.show()
```




```
[62]: for val in [False, True]:
        tmp = reg_df[reg_df["isStrength"] == val]
        plt.scatter(tmp["heartRateQ99"],
                    tmp["kiloCalories"],
                    s=3,
                    label=val)

plt.xlabel("0.99 quantile of heart rate (bpm)")
plt.ylabel("Kilocalories")
plt.legend(title="isStrength", loc="best")
plt.savefig("./img/99q_hr_vs_kilocalories_scatter_by_strength.png")
plt.show()
```



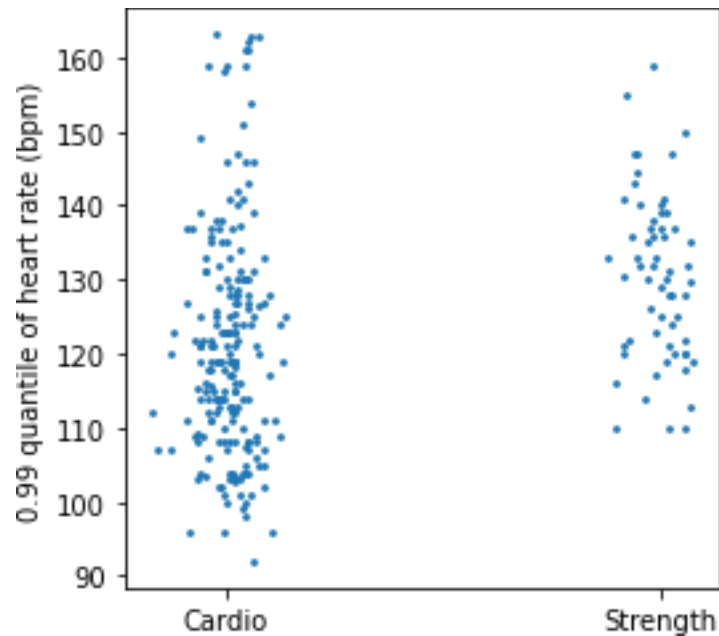
```
[63]: plt.scatter(reg_df["isStrength"] + np.random.normal(scale=1/20, _
    size=len(reg_df)),
    reg_df["kiloCalories"], s=3)

plt.ylabel("Kilocalories")
plt.xticks(ticks=[0, 1], labels=["Cardio", "Strength"])
plt.savefig("./img/is_strength-vs_kilocalories_jitter.png")
plt.show()
```




```
[64]: plt.scatter(reg_df["isStrength"] + np.random.normal(scale=1/20, _
    ↪ size=len(reg_df)),
    reg_df["heartRateQ99"], s=3)

plt.ylabel("0.99 quantile of heart rate (bpm)")
plt.xticks(ticks=[0, 1], labels=["Cardio", "Strength"])
plt.savefig("./img/is_strength_vs_99q_hr_scatter.png")
plt.show()
```



```
[65]: plt.scatter(reg_df["isStrength"] + np.random.normal(scale=1/20, _  
    ↪ size=len(reg_df)),  
        reg_df["totalTime"], s=3)  
  
plt.ylabel("Time (minutes)")  
plt.xticks(ticks=[0, 1], labels=["Cardio", "Strength"])  
plt.savefig("./img/is_strength_vs_time_jitter.png")  
plt.show()
```

8.4 Correlation

We convert binary the feature isStrength to integers for the rest of the analysis.

```
[66]: reg_df["isStrength"] = reg_df["isStrength"].astype(int)
```

We inspect the correlation matrix to check for multicollinearity. It should be noted that the correlation between kiloCalories and totalTime is quite high and this to be expected.

```
[67]: C = reg_df.corr(method='pearson')
# C = C.style.background_gradient(cmap='YlGn')
# C = C.set_precision(2)
C
```

```
[67]:
```

	kiloCalories	totalTime	heartRateQ99	isStrength
kiloCalories	1.00	0.92	0.51	0.55
totalTime	0.92	1.00	0.28	0.69
heartRateQ99	0.51	0.28	1.00	0.26
isStrength	0.55	0.69	0.26	1.00

8.5 Multicollinearity

We inspect the respect variance inflation factors and are happy to see that all are below 10.

```
[68]: tmp = reg_df.drop(["kiloCalories", "sport"], axis=1)
```

```

vifs = []
for i in range(tmp.shape[1]):
    vif = variance_inflation_factor(tmp.to_numpy(), i)
    vifs.append(round(vif, 2))

vifs = pd.DataFrame(vifs, index=tmp.columns, columns=['VIF'])
vifs = vifs.sort_values('VIF', ascending=False)
vifs = vifs.reset_index()
vifs = vifs.rename(columns={'index': 'Variable'})

# vifs = vifs.style.background_gradient(cmap='OrRd')
# vifs = vifs.set_precision(2)

vifs

```

```

[68]:      Variable  VIF
0      totalTime  6.30
1    heartRateQ99  3.90
2      isStrength  2.41

```

8.6 Modelling

Before the actual modelling we prepare a function to calculate RMSE to compare models and extract the true kiloCalories into a separate array.

```

[69]: y_true = reg_df['kiloCalories'].to_numpy()

```

```

[70]: def calc_rmse(y_true, y_pred):
      x = np.sqrt(np.mean(np.power(y_true - y_pred, 2)))
      return round(x, 4)

```

```

[71]: all_results = []

```

8.6.1 Time only

We start the modelling section of by building the simplest model that comes to mind: predict kiloCalories using totalTime.

```

[72]: formula = "kiloCalories ~ totalTime"
      mdl_time = smf.ols(formula=formula, data=reg_df)
      mdl_time = mdl_time.fit()
      mdl_time.summary2()

```

```

[72]: <class 'statsmodels.iolib.summary2.Summary'>
      """

```

Results: Ordinary least squares

=====

Model:	OLS	Adj. R-squared:	0.849
Dependent Variable:	kiloCalories	AIC:	3070.3883
Date:	2020-10-06 10:25	BIC:	3077.5477
No. Observations:	265	Log-Likelihood:	-1533.2
Df Model:	1	F-statistic:	1483.
Df Residuals:	263	Prob (F-statistic):	4.24e-110
R-squared:	0.849	Scale:	6254.0

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
Intercept	17.1713	8.7686	1.9583	0.0513	-0.0942	34.4369
totalTime	6.7722	0.1759	38.5076	0.0000	6.4260	7.1185

Omnibus:	35.699	Durbin-Watson:	1.803
Prob(Omnibus):	0.000	Jarque-Bera (JB):	92.425
Skew:	0.605	Prob(JB):	0.000
Kurtosis:	5.628	Condition No.:	90

=====

.....

```
[73]: y_pred = mdl_time.predict(reg_df)
      rmse = calc_rmse(y_pred, y_true)
      all_results.append((rmse, formula))
```

```
[74]: print(rmse)
```

78.7832

8.6.2 By sport

The next regression we are going to do will be univariate regression separately for each sport, this will help us answer the question which sport is the most effective at burning calories during a workout.

```
[75]: all_sports = sorted(reg_df["sport"].unique())
      reg_sports_res = []

      # For all sport do simple linear regression
      for sport in all_sports:
          tmp = reg_df[reg_df["sport"] == sport]
          formula = 'kiloCalories ~ totalTime'
          mdl_sport = smf.ols(formula=formula, data=tmp)
          mdl_sport = mdl_sport.fit()
          sport_stats = [formula, sport] + list(mdl_sport.params) + [mdl_sport.
↪rsquared]
          reg_sports_res.append(sport_stats)
```

```
cols = ["Formula", "Sport", "Intercept", "Slope", "R squared"]

reg_sports_res = pd.DataFrame(reg_sports_res, columns=cols)
reg_sports_res = reg_sports_res.sort_values(["Slope"], ascending=False)
reg_sports_res = reg_sports_res.reset_index(drop=True)

readme_df = reg_sports_res.copy().round(2)

# reg_sports_res = reg_sports_res.style.background_gradient(cmap='YlGn',
# subset='Slope')
# reg_sports_res = reg_sports_res.set_precision(2)

reg_sports_res
```

```
[75]:
```

	Formula	Sport	Intercept	Slope	R squared
0	kiloCalories ~ totalTime	treadmill_running	-21.23	10.14	0.96
1	kiloCalories ~ totalTime	cycling	-9.73	7.44	0.98
2	kiloCalories ~ totalTime	walking	12.59	6.95	0.82
3	kiloCalories ~ totalTime	strength_training	-12.73	6.76	0.44

8.7 Time and heart rate

We try to enhance the model by adding heartRateQ99.

```
[76]: formula = "kiloCalories ~ totalTime + heartRateQ99"
mdl_time_and_hr = smf.ols(formula=formula, data=reg_df)
mdl_time_and_hr = mdl_time_and_hr.fit()
mdl_time_and_hr.summary2()
```

```
[76]: <class 'statsmodels.iolib.summary2.Summary'>
      .....
```

Results: Ordinary least squares

```
=====
Model:                OLS                Adj. R-squared:    0.918
Dependent Variable:    kiloCalories        AIC:                2910.3348
Date:                  2020-10-06 10:25    BIC:                2921.0740
No. Observations:      265                Log-Likelihood:      -1452.2
Df Model:               2                  F-statistic:         1472.
Df Residuals:           262                Prob (F-statistic):   3.32e-143
R-squared:              0.918              Scale:              3405.9
=====
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
Intercept	-420.7741	30.1664	-13.9484	0.0000	-480.1736	-361.3746
totalTime	6.2025	0.1353	45.8339	0.0000	5.9361	6.4690
heartRateQ99	3.7435	0.2519	14.8636	0.0000	3.2476	4.2394

```
=====
```

Omnibus:	14.259	Durbin-Watson:	1.917
Prob(Omnibus):	0.001	Jarque-Bera (JB):	21.717
Skew:	-0.355	Prob(JB):	0.000
Kurtosis:	4.210	Condition No.:	1106

=====

* The condition number is large (1e+03). This might indicate strong multicollinearity or other numerical problems.

```
[77]: y_pred = mdl_time_and_hr.predict(reg_df)
      rmse = calc_rmse(y_pred, y_true)
      all_results.append((rmse, formula))
```

```
[78]: print(rmse)
```

58.0287

8.8 Time with random effects by workout type

```
[79]: formula = 'kiloCalories ~ totalTime + heartRateQ99'
      re_formula = '~ totalTime'
      group = 'isStrength'

      mdl_time_with_hr_re = smf.mixedlm(formula=formula,
                                         data=reg_df,
                                         groups=reg_df[group],
                                         re_formula=re_formula)

      mdl_time_with_hr_re = mdl_time_with_hr_re.fit(method='lbfgs')
      mdl_time_with_hr_re.summary()
```

```
[79]: <class 'statsmodels.iolib.summary2.Summary'>
      *****
```

Mixed Linear Model Regression Results

```
=====
Model:                MixedLM    Dependent Variable:    kiloCalories
No. Observations:      265        Method:                REML
No. Groups:            2          Scale:                2570.3031
Min. group size:       61         Log-Likelihood:       -1416.8825
Max. group size:       204        Converged:            Yes
Mean group size:       132.5
=====
```

	Coef.	Std.Err.	z	P> z	[0.025	0.975]
Intercept	-475.991	65.174	-7.303	0.000	-603.730	-348.252
totalTime	6.755	0.858	7.874	0.000	5.074	8.437
heartRateQ99	3.933	0.220	17.868	0.000	3.502	4.365

```

Group Var          210.604  51.671
Group x totalTime Cov   9.904   0.550
totalTime Var          0.466   0.036
=====

.....

```

```
[80]: y_pred = mdl_time_with_hr_re.predict(reg_df)
      rmse = calc_rmse(y_pred, y_true)
      all_results.append((rmse, formula, re_formula, group))
```

```
[81]: print(rmse)
```

```
60.9247
```

8.9 Model evaluation

We compare the linear models created earlier:

```
[82]: comp_df = pd.DataFrame(all_results, columns=["RMSE", "Formula", "Random_
      ↪effects", "Groups"])
      comp_df = comp_df.sort_values("RMSE")

      # comp_df = comp_df.style.background_gradient(cmap='OrRd', subset='RMSE')
      # comp_df = comp_df.set_precision(2)

      comp_df
```

```
[82]:   RMSE          Formula Random effects    Groups
1  58.03 kiloCalories ~ totalTime + heartRateQ99      None      None
2  60.92 kiloCalories ~ totalTime + heartRateQ99 ~ totalTime isStrength
0  78.78          kiloCalories ~ totalTime      None      None
```


For further evaluation we choose the random effects model.

```
[83]: mdl = mdl_time_with_hr_re
      residuals = mdl_time_with_hr_re.resid
```

8.9.1 Visual inspection

We proceed to inspect the residuals of the model. First we view the histogram of the residuals. It can be seen that it looks normal.

```
[84]: plt.hist(residuals)
      plt.ylabel("Frequency")
      plt.xlabel("Residuals")
      plt.savefig("./img/mdl_residuals.png")
      plt.show()
```




The next plot is a qqplot created to visually inspect the normality of the residuals. We see 3 nasty outliers in the top right corner.

```
[85]: plt.figure() ax
      = plt.gca()


      qqplot(data=mdl.resid,
              ax=ax,
              color='#1f77b4',
              markersize=3,
              line='45',
              fit=True,
              alpha=1/2)

      plt.savefig('./img/mdl_qq.png')
      plt.show()
```



The third plot we make is a plot of the standardized residuals to check for homoskedasticity. Again we see the same outliers as on the plot above.

```
[86]: residuals_std = np.abs((residuals - np.mean(residuals)) / np.std(residuals))
plt.plot(residuals_std, "o", markersize=2)
plt.xlabel("Observation")
plt.ylabel("Standardized residuals")
plt.savefig("./img/mdl_residuals_std.png")
plt.show()
```





Finally we compare the predicted kiloCalories with the actual values.

```
[87]: y_pred = mdl.predict(reg_df)
      y_pred = y_pred.to_numpy().reshape(len(y_pred))

      m = np.min(np.hstack([y_true, y_pred]))
      M = np.max(np.hstack([y_true, y_pred]))

      x = np.linspace(m, M, len(y_pred))
      plt.plot(y_true, y_pred, 'o', markersize=2)
      plt.plot(x,x, alpha=3/4)
      plt.ylabel('Predicted')
      plt.xlabel('Actual')
      plt.tight_layout()
      plt.savefig('./img/mdl_predicted_vs_actual.png')
      plt.show()
```



The next step is to take a look at the data points with the biggest error. As can be seen the model has issues predicting strength training workouts.

```
[88]: errors = reg_df.copy() errors['kiloCaloriesPredicted']
      = mdl_predict(reg_df)

errors['error'] = np.abs(errors['kiloCalories'] -
      errors['kiloCaloriesPredicted'])

errors = errors.sort_values('error', ascending=False)
errors = errors.reset_index(drop=True)

order = ['kiloCaloriesPredicted',
        'kiloCalories',
        'error',
        'totalTime',
        'isStrength']

errors = errors[order]

errors = errors.head(5)

errors = errors.style.background_gradient(cmap='OrRd', subset='error')
```

```
errors = errors.set_precision(2)

errors
```

[88]: <pandas.io.formats.style.Styler at 0x7f8538793640>

9 Summary

```
[89]: # Make table for README
# print(tabulate.tabulate(by_sport.values, by_sport.columns, tablefmt="pipe"))
```

```
[90]: # Make table for README
# print(tabulate.tabulate(readme_df.values, readme_df.columns, tablefmt="pipe"))
```

- In this project I define a workout as each instance in time when my watch was recording me.
- I downloaded data generated by my Polar watch that tracks heart rate and estimates burned kilocalories during workouts.
- The data came in the form of .json files which were read, transformed and cleaned with pandas.
- The clean dataset contains 283 workouts over a nearly one year period during which I burned roughly 12kg of body fat.

Sport	Total kilocalories	Total kilograms
walking	33080	4.3
strength_training	31547	4.1
treadmill_running	19825	2.57
cycling	4029	0.52
running	940	0.12

- The timing of my workouts appears to follow a bimodal distribution with peaks at 12:00 and 20:00.
- After further transforming the data, I find that the duration of a workout and kilocalorie's burned have a 0.92 correlation.
- Several linear regressions were performed.
- kilocalories ~ duration on the entire dataset achieved $R^2 = 0.85$ and $RMSE = 79$.
- Regressions were performed on subsets of the data, specifically by sport - the highest slope is 10.14 kiloCalories per minute.

Formula	Sport	Intercept	Slope	R squared
kilo_calories ~ total_time	treadmill_running	-21.23	10.14	0.96
kilo_calories ~ total_time	cycling	-9.73	7.44	0.98
kilo_calories ~ total_time	walking	12.59	6.95	0.82

Formula	Sport	Intercept	Slope	R squared
kilo_calories ~ total_time	strength_training	-12.73	6.76	0.44

- A linear mixed model with random effects was created and validated. It achieved a RMSE = 61 and normal looking residuals.
- The biggest errors made by the mixed model was on strength training data points.